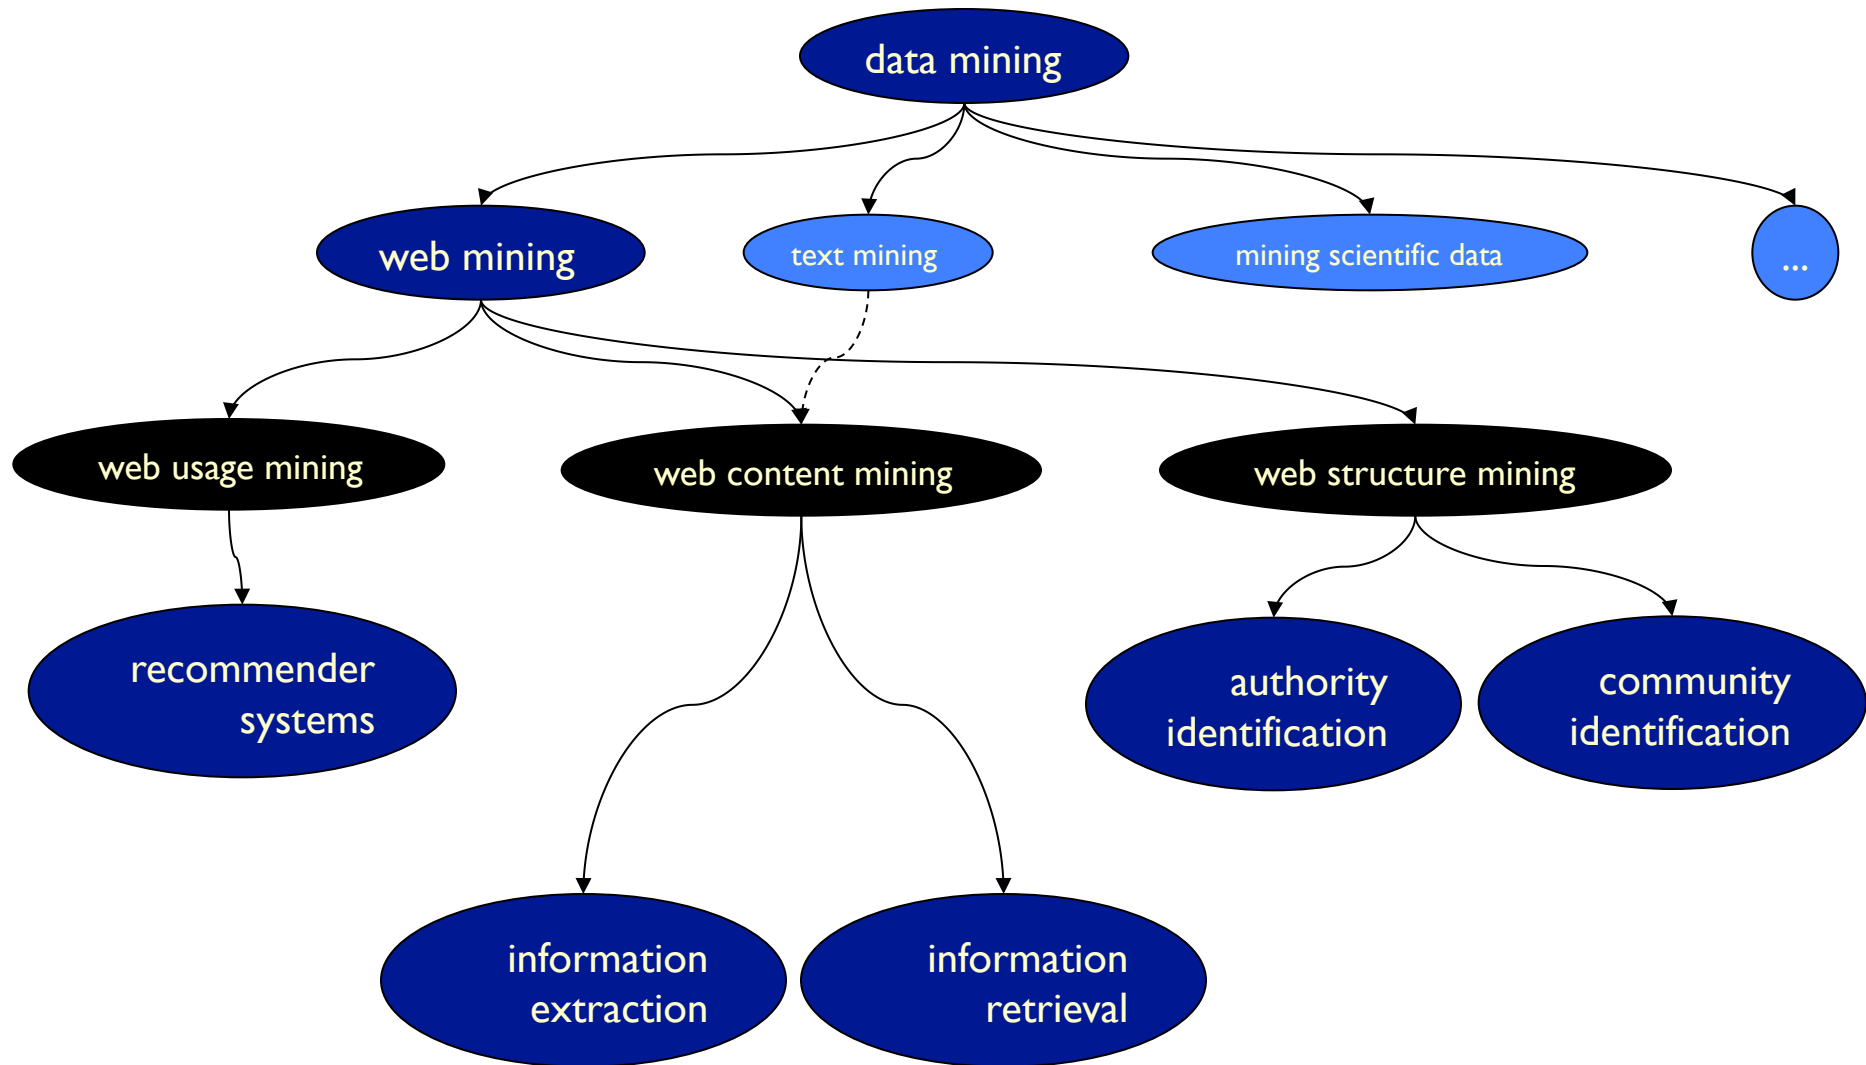


Web Mining: Structure: Link Analysis

Alípio Jorge, DCC-FC, Universidade do Porto
amjorge@fc.up.pt



Web Mining

▶ Web Usage Mining

- ▶ discovery from user access patterns from logs or alike
- ▶ applications:
 - ▶ user segmentation, recommendation, personalization, adaptation, usability improvement

▶ Web Structure Mining

- ▶ discovery of useful knowledge from hyperlinks
- ▶ applications:
 - ▶ discover important pages (information retrieval)
 - ▶ discover communities

▶ Web Content Mining

- ▶ extracts information from Web pages
- ▶ applications
 - ▶ information extraction, summarization, topic extraction, discovering user emotions

Web structure mining

- ▶ Take advantage of the information in web hyperlinks
 - ▶ links are created locally
 - ▶ web structure, as a whole, is not planned.
- ▶ Take advantage of the information in social links
 - ▶ social networks
- ▶ To understand the structure of the web
 - ▶ Link analysis
 - ▶ Analysis of the topology of connections

Web graphs

- ▶ Internet can be seen as different interdependent graphs
 - ▶ pages and hyperlinks (web)
 - ▶ computers and communications between them (internet)
- ▶ Web graph
 - ▶ very large (2×10^{10} ?)
 - ▶ dynamical (changes structure and content)
 - ▶ has virtual parts (dynamic pages, harder to analyse)
 - ▶ disconnected (has islands)
 - ▶ sparse (relatively few connections)

How can we use web structure?


- ▶ For finding
 - ▶ prestigious web pages
 - ▶ central links in social webs
 - ▶ communities
 - ▶ web page clusters pointing to each other
 - ▶ groups of people who change emails
- ▶ Studying web structure is related to
 - ▶ social network analysis
 - ▶ e.g. package "sna" of R
 - ▶ complex networks

Network analysis

► Interesting phenomena in a network


► central nodes

- are important to connect two parts of the network
- are involved in many indirect connections



► prestigious nodes


- tend to be referred to by many other nodes



Network analysis: centrality measures

- ▶ Degree centrality of a node
 - ▶ network has n nodes (actors)
 - ▶ $d(i)$ is the number of links of node i – node degree
 - ▶ the more links, the higher centrality
 - ▶ range $[0, 1]$
- ▶ Undirected graph


$$C_D(i) = \frac{d(i)}{n-1}$$



Network analysis: centrality measures

- ▶ Degree centrality of a node
 - ▶ network has n nodes (actors)
 - ▶ $d_o(i)$ is the number of **out-links** of node i – out-degree
 - ▶ the more links, the higher centrality
 - ▶ range $[0,1]$
- ▶ Directed graph

$$C'_D(i) = \frac{d_o(i)}{n-1}$$




Network analysis: centrality measures

► Closeness centrality of a node

- a node is important if it is closer to all other nodes
- $d(i,j)$ is the distance between nodes i and j – e.g. number of edges
- range $[0,1]$ (assuming a connected graph)

► Undirected graph

$$C_D(i) = \frac{n-1}{\sum_{j=1}^n d(i,j)}$$




Network analysis: centrality measures

► Closeness centrality of a node

- a node is important if it is closer to all other nodes
- $d(i,j)$ is the distance between nodes i and j – e.g. number of edges
- range $[0,1]$ (assuming a connected graph)

- Directed graph – distance now considers direction

$$C_D(i) = \frac{n-1}{\sum_{j=1}^n d(i,j)}$$




Network analysis: centrality measures

► Betweenness centrality of a node


- a node is important if it is between other nodes
- p_{jk} is the number of shortest paths between j and k
- $p_{jk}(i)$ is the number of shortest paths between j and k that go through i ($i \neq j, i \neq k$)
 - range $[0, (n-1)(n-2)/2]$
- Undirected graph

$$C_B(i) = \sum_{j < k, j \neq i, k \neq i} \frac{p_{jk}(i)}{p_{jk}}$$



Network analysis: centrality measures

- ▶ There is data about friendship requests in a social net
 - ▶ Who would you pick as a marketing mate: f or d?
 - ▶ Who would you pick for collecting information?
 - ▶ Who would you pick for distribution of goods?



Network analysis: prestige measures


- ▶ Degree prestige of a node

- ▶ a node is prestigious if it is referred by other nodes

- ▶ Directed graph

- ▶ $d_i(i)$ is the number of **in-links** of node i – in-degree

$$P_D(i) = \frac{d_i(i)}{n-1}$$



Network analysis: prestige measures

- ▶ node A is referred by n ordinary nodes
- ▶ node B is referred by n nodes, k of which prestigious
 - ▶ which node has higher prestige?
- ▶ we must take the prestige of pointing nodes into account
- ▶ HITS and PageRank do just that


Using web structure for information retrieval

▶ Search

- ▶ Search a page about topic X
- ▶ Each page Y is relevant according to
 - ▶ similarity between the content of X and Y

▶ Link analysis

- ▶ Each page Y is relevant according to
 - ▶ number of references to page Y
 - ▶ content of pages which refer to Y




- ▶ Pages linked to pages with interesting content are also potentially interesting

HITS (hyperlink induced topic search)

[Chakrabarti et al.]

- ▶ Discovery of two kinds of pages
 - ▶ **Authorities**
 - ▶ pages referred to by many others in a specific topic
 - ▶ **Hubs**
 - ▶ pages that refer to many others
- ▶ In a first stage we use text similarity then we use link structure



Hubs e Authorities

- ▶ **Relevance of an Authority**

- ▶ if a page is referred to by many others, then it must be relevant
- ▶ it enables search more robust to variation in terms
 - ▶ example “data mining” and “machine learning”

- ▶ **Quality of a Hub**

- ▶ if a hub refers to many important authorities then it is a good hub

- ▶ **The relevance of an Authority and the quality of a Hub are interdependent**

Link analysis with HITS

- ▶ Community discovery about a topic by computing hubs and authorities to that topic
 1. given a query (topic) Q , collect a set of seed pages $S = \{s_1, s_2, \dots, s_n\}$ (root set)
 2. S is expanded to $T = S \cup \{d \mid s \rightarrow d \text{ or } d \rightarrow s, s \in S\}$
 3. initially, each page $r \in T$ has authority weight $\mathbf{a}(\mathbf{r}) = 1$, hub weight $\mathbf{h}(\mathbf{r}) = 1$

$$a(r) = \sum_{d \rightarrow r} h(d)$$


For each page we update the values of \mathbf{a} and \mathbf{h}

$$h(r) = \sum_{r \rightarrow d} a(d)$$

Normalize \mathbf{a} and \mathbf{h} and repeat step 3 until convergence (typically 10 it.)

4. The community corresponds to the k top pages with highest \mathbf{a} and \mathbf{h}

Link analysis



$x1 \rightarrow s1$

$x2 \rightarrow s1$

$s1 \rightarrow y1$

$s2 \rightarrow x1$

$x2 \rightarrow s3$

$s3 \rightarrow y1$

$a(s1)=1, h(s1)=1$

$a(s2)=1, h(s2)=1$

$a(s3)=1, h(s3)=1$

$a(x1)=1, h(x1)=1$

$a(x2)=1, h(x2)=1$

$a(y1)=1, h(y1)=1$

iteração 1

$a1(s1)=2, h1(s1)=1$

$a1(s2)=0, h1(s2)=1$

$a1(s3)=1, h1(s3)=1$

$a1(x1)=1, h1(x1)=1$

$a1(x2)=0, h1(x2)=2$

$a1(y1)=2, h1(y1)=0$

iteração 1 (norm.)

$a2(s1)=1, h2(s1)=0,5$

$a2(s2)=0, h2(s2)=0,5$

$a2(s3)=0,5, h2(s3)=0,5$

$a2(x1)=0,5, h2(x1)=0,5$

$a2(x2)=0, h2(x2)=1$

$a2(y1)=1, h2(y1)=0$

iteração 2

$a2(s1)=1,5, h2(s1)=1$

$a2(s2)=0, h2(s2)=0,5$

$a2(s3)=1, h2(s3)=1$

$a2(x1)=0,5, h2(x1)=1$

$a2(x2)=0, h2(x2)=1,5$

$a2(y1)=1, h2(y1)=0$

iteração 2 (norm.)

$a2(s1)=1, h2(s1)=0,66$

$a2(s2)=0, h2(s2)=0,33$

$a2(s3)=0,66, h2(s3)=0,66$

$a2(x1)=0,33, h2(x1)=0,66$

$a2(x2)=0, h2(x2)=1$

$a2(y1)=0,66, h2(y1)=0$

iteração 3

$a2(s1)=1,66, h2(s1)=0,66$

$a2(s2)=0, h2(s2)=0,33$

$a2(s3)=1, h2(s3)=0,66$

$a2(x1)=0,33, h2(x1)=1$

$a2(x2)=0, h2(x2)=1,66$

$a2(y1)=1,33, h2(y1)=0$

iteração 3 (norm.)

$a2(s1)=1, h2(s1)=0,4$

$a2(s2)=0, h2(s2)=0,2$

$a2(s3)=0,6, h2(s3)=0,4$

$a2(x1)=0,2, h2(x1)=0,6$

$a2(x2)=0, h2(x2)=1$

$a2(y1)=0,8, h2(y1)=0$

topo authority: $s1 (1), y1 (0,8), s3 (0,6)$

topo hub: $x2 (1), x1 (0,6), s1 (0,4)$

HITS with an adjacency matrix

- ▶ The graph of connections / links can be represented by an adjacency matrix A

- ▶ Where

- ▶ a not normalized (ann) is

$$\text{ann} = A^T h$$

- ▶ a normalized (a) is

$$a = \text{ann} / \max(\text{ann})$$

- ▶ similarly

$$\text{hnn} = A.a; \quad h = \text{hnn} / \max(\text{hnn})$$

	[s1]	[s2]	[s3]	[x1]	[x2]	[y1]
[s1]	0	0	0	0	0	1
[s2]	0	0	0	1	0	0
[s3]	0	0	0	0	0	1
[x1]	1	0	0	0	0	0
[x2]	1	0	1	0	0	0
[y1]	0	0	0	0	0	0

HITS in pseudo-code

- ▶ Graph of connections given as an adjacency matrix
- ▶ Given a number of iterations

hits-iterate(A)

$a_0 \leftarrow h_0 \leftarrow (1, 1, \dots, 1)$

$k \leftarrow 1$

Repeat

$hnn_k \leftarrow A \cdot a_{k-1}$

$a_k \leftarrow ann_k / \max(ann_k)$

$h_k \leftarrow hnn_k / \max(hnn_k)$

Until $|a_k - a_{k-1}| < ea$ and $|h_k - h_{k-1}| < eh$

return ann_k , a_k , hnn_k , h_k

```
> A
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]    0    0    0    0    0    1
[2,]    0    0    0    1    0    0
[3,]    0    0    0    0    0    1
[4,]    1    0    0    0    0    0
[5,]    1    0    1    0    0    0
[6,]    0    0    0    0    0    0
```

```
> hits(A,4)
      [,1] [,2] [,3] [,4]
[1,]  1.6  0.8 1.000 0.500
[2,]  0.0  0.2 0.000 0.125
[3,]  1.0  0.8 0.625 0.500
[4,]  0.2  1.0 0.125 0.625
[5,]  0.0  1.6 0.000 1.000
[6,]  0.8  0.0 0.500 0.000
```

22 amjorge@fc.up.pt

HITS by eigenvectors

$$a = A^T h \qquad h = Aa$$

$$a = A^T Aa \qquad h = AA^T h$$

- ▶ with normalization

$$a = A^T .k_1 h \qquad h = A.k_2 a$$

$$a = A^T A.k_2 a \qquad h = AA^T .k_1 h$$

- ▶ a is the largest eigenvector of $A^T A$
- ▶ h is the largest eigenvector of $A.A^T$
- ▶ x eigenvector of M if $Mx=k.x$, where k is a scalar

And on R we get

```
> t(eigen(t(A)%*%A)$vectors[,1])  
      [,1] [,2]      [,3] [,4] [,5] [,6]  
[1,] 0.8506508      0 0.5257311      0      0      0
```

- ▶ i.e., authorities are s1 and s3

```
> t(eigen(A%*%t(A))$vectors[,1])  
      [,1] [,2] [,3]      [,4]      [,5] [,6]  
[1,]      0      0      0 0.5257311 0.8506508      0
```

- ▶ hubs are x2 and x1
- ▶ Function eigen computes eigenvectors (among other things)

Comments on eigenvectors

- ▶ Iterative algorithm finds the principal eigenvectors
 - ▶ Major communities
- ▶ Other eigenvectors
 - ▶ Alternative communities
 - ▶ E.g. query “classification” or “football”
- ▶ Convergence
 - ▶ HITS always converges
 - ▶ Different initializations may give different results
 - ▶ If there are repeated principal eigenvalues
 - ▶ If $A^T A$ is reducible
 - the graph is not strongly connected


Notes

- ▶ Search in HITS starts by CONTENT relevance
 - ▶ root set
- ▶ then content is ignored
 - ▶ only links are exploited
- ▶ Example:
 - ▶ look for pages of "japanese car manufacturers"
 - ▶ the page of Honda will not have this page
 - ▶ software companies...
 - ▶ pages are not typically self descriptive

How to identify inlinks

- ▶ **google.com**
 - ▶ Query “link: <url>”
 - ▶ no space after :
- ▶ results are a sample of the actual set of links

Activity





- ▶ Determine the most interesting hub
- ▶ Determine the most important authority
- ▶ Suppose we are looking for information about a car model X and page A contains that model, how would that change your previous results?
- ▶ Think of way for enhancing recommendation using hits

PageRank

- ▶ HITS was proposed in January 1998
 - ▶ Kleinberg
- ▶ PageRank was proposed in April 1998 and is used by Google
 - ▶ Sergey Brin and Larry Page
- ▶ HITS and PageRank have many similarities
- ▶ but they have very important differences
 - ▶ computational
 - ▶ robustness of results
- ▶ The idea of PageRank
 - ▶ rank pages according to their prestige
 - ▶ prestige is (mainly) determined by **inlinks** and their prestige

PageRank: The idea

- ▶ We consider a random robot
 - ▶ $\text{prob}(\text{página B} \rightarrow \text{página A}) = 1/n$
 - ▶ n is the number of outlinks from B
 - ▶ $\text{prob}(\text{being on page A coming from B}) = \text{prob(B)}/n$
 - ▶ $\text{prob(A)} =$
 $\text{prob(B)}/\text{Out(B)} + \text{prob(B2)}/\text{Out(B2)} + \text{prob(B3)}/\text{Out(B3)}$
- ▶ Most important pages will have higher probability
- ▶ What about loops and direct accesses?



PageRank Mathematics

- ▶ $R(i) = R(j_1) / O_{j_1} + R(j_2) / O_{j_2} + \dots$
- ▶ How can we determine $R(i)$?
 - ▶ system of n equations and n unknowns
- ▶ $R = \langle R(1), R(2), \dots, R(n) \rangle$
- ▶ $A_{ij} = 1/O_i$ OR zero
- ▶ $R = A^T \cdot R$
- ▶ This could be enough, but...

> Adjacency matrix

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	0	0	0	0	0	1
[2,]	0	0	0	1	0	0
[3,]	0	0	0	0	0	1
[4,]	1	0	0	0	0	0
[5,]	1	0	1	0	0	0
[6,]	0	0	0	0	0	0

> A

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	0	0	0	0	0	1
[2,]	0	0	0	1	0	0
[3,]	0	0	0	0	0	1
[4,]	1	0	0	0	0	0
[5,]	0,5	0	0,5	0	0	0
[6,]	0	0	0	0	0	0

PageRank – some problems


- ▶ $R = A^T.R$
- ▶ For the above to have a unique solution A must be
 - ▶ **stochastic** (all rows must sum 1)
 - ▶ often it is not: there are nodes with no outlinks
 - ▶ solution 1: remove nodes without outlinks
 - ▶ solution 2: artificially insert equal weights into a row with zeros

```
> A
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,] 0    0    0    0    0    1
[2,] 0    0    0    1    0    0
[3,] 0    0    0    0    0    1
[4,] 1    0    0    0    0    0
[5,] 0,5  0    0,5  0    0    0
[6,] 0    0    0    0    0    0
```

```
> A
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,] 0    0    0    0    0    1
[2,] 0    0    0    1    0    0
[3,] 0    0    0    0    0    1
[4,] 1    0    0    0    0    0
[5,] 0,5  0    0,5  0    0    0
[6,] 1/6  1/6  1/6  1/6  1/6  1/6
```


PageRank – some problems

- ▶ $R = A^T.R$
- ▶ ...A must be
 - ▶ **irreducible** (in the graph there is a path from any node to any other node)
 - ▶ often it is not the case (there is no path from S1 to S2)
 - ▶ **aperiodic** (the greatest common divisor of all cycles for each node is 1)
 - ▶ $A \rightarrow B, B \rightarrow C, C \rightarrow A$: the cycle has period 3
 - ▶ No loop traps
- ▶ **Solution**
 - ▶ add a link to every two pages
 - ▶ in fact, if one is in one page can go directly to any other by typing its URL
 - ▶ the probability of transition is controlled by a parameter d



PageRank: Teleportation

- ▶ When on a page B, there is a certain probability (ex: 0.1) of teleporting to a page A which has no direct connection to B
 - ▶ $\text{prob}(\text{getting to A}) = 0.1 / (\text{number of nodes}) + 0.9 * \text{prob}(\text{direct access})$
- ▶ $R(A)$ is proportional to this probability


PageRank

<http://www-db.stanford.edu/~backrub/google.html>

<http://www.iprcom.com/papers/pagerank/>

$$R(A) = (1-d) + d \cdot (R(T1)/Out(T1) + \dots + R(Tn)/Out(Tn))$$

- ▶ $R(X)$ page rank of page X
- ▶ d damp factor (solves connectivity problems and models direct accesses)
- ▶ $Out(X)$ number of outlinks of X



PageRank additional criteria*

- ▶ *(improving the user model)*
- ▶ Visibility of a link
- ▶ Position of a link within a document
- ▶ Distance between web pages
 - ▶ same server, same domain, same region
- ▶ Importance of a linking page
- ▶ Up-to-dateness of a linking page

*<http://pr.efactory.de/e-further-factors.shtml>

Algorithm

- ▶ We can solve a system of equations
- ▶ We can calculate R iteratively
 - ▶ assign initial R values to pages
 - ▶ calculate new values for R
 - ▶ iterate (number of iterations depends on the size of the network)


PageRank and HITS

- ▶ PageRank can be computed offline
 - ▶ it is query independent
 - ▶ which can be a disadvantage: a page can be an authority in a topic but not in general
 - comparar www.publico.pt ou www.oftalmologia.pt
- ▶ PageRank is more robust to SPAM
 - ▶ importance of a page depends on inlinks not outlinks
- ▶ PageRank does not consider time
- ▶ PageRank is more robust to perturbations in the input than HITS

Web Spamming

- ▶ Artificially increasing the rank of a page without increasing its specific information value
 - ▶ Search Engine Optimization can be spam or not
 - ▶ debateable (<http://www.webworkshop.net/ethical-search-engine-optimization.html>)
- ▶ Content Spamming
 - ▶ insert popular words (even if unrelated)
 - ▶ repeating important terms
 - ▶ dumping many unrelated terms
- ▶ Link Spamming
 - ▶ outlink spamming: directory cloning
 - ▶ inlink spamming
 - ▶ honey-pot
 - ▶ submit URLs to Web Directories
 - ▶ Posting links to forums or the like
 - ▶ link exchange schemes
 - ▶ spam farms

Activity



- ▶ Assume damp factor of 0.9
- ▶ Suppose the PageRank of A is 1, what is the PR of B?
- ▶ and of C?
- ▶ Determine the PageRank of the pages of the graph

Community Discovery

- ▶ **Community:** group of entities (people, organizations) sharing common interests.
 - ▶ Users who like metal music
 - ▶ Treckies
- ▶ **What for?**
 - ▶ Source of resources for users with similar interests
 - ▶ Sociology of the web: we know better, we can exploit better
 - ▶ Target advertising

Community Discovery

- ▶ Given a set of entities S
 - ▶ Of the same type
- ▶ A community is
 - ▶ A pair $C = (\text{Theme}, \text{Group})$
- ▶ Example
 - ▶ Users who like metal music

Communities

▶ Web pages

- ▶ Users in the same community are usually interconnected through hyperlinks
- ▶ Pages contain words that reveal the theme

▶ Emails


- ▶ Members of a community exchange emails (links)
- ▶ Emails contain words revealing the theme

▶ Documents

- ▶ Members of a community are more likely to appear together in the same sentences or documents (this is the link)
- ▶ Words indicate the community theme

Algorithm: Bipartite Core Communities


- ▶ **Bipartite graph of Fans and Centers**
 - ▶ Music fans and band pages
- ▶ **Identifying (i,j) cores**
 - ▶ i fans and j centers
 - ▶ Fans ~ Hubs, Centers ~ Authorities
- ▶ **We could use HITS**
 - ▶ But computing eigenvalues is relatively inefficient
 - ▶ We will describe an algorithm by R. Kumar



Algorithm: Bipartite Core Communities

► Pruning

- Delete pages that are too highly referenced
 - $\text{Inlink} > 500$
- Prune fans and centers
 - Fans with $\text{outdegree} < j$
 - Centers with $\text{indegree} < i$
 - Example for $(i=3, j=2)$



After pruning


Algorithm: Generating all (i,j) cores

► After Pruning

- Fix j , start with all $(1,j)$ cores
 - Set of fans with outdegree at least j
- Look for $(2,j)$ cores by checking every fan that points to a center in a $(1,j)$ core
- Similarly for $(3,j)$ in a APRIORI fashion

► Note

- This algorithm finds cores of communities, not the whole community




Suppose this was the result of pruning

Algorithm: Generating all (i,j) cores

► Example

- Find all (3,2) cores
 - Identify (1,2) cores
 - Fans with 2 outlinks (min)
 - Identify (2,2) cores
 - Combine pairs of fans to find larger cores
 - Identify (3,2) cores
- Find all (3,3) cores



Suppose this was the result of pruning

Resources

- ▶ Books

- ▶ Web Data Mining, Bing Liu, Springer, 2007
- ▶ Mining the World Wide Web, Chang, G., Healey, M., McHugh, J., Wang, J., Kluwer Academic Press, 2001.

- ▶ Article

- ▶ Google's PageRank Explained and how to make the most of it, Phil Craven, <http://www.webworkshop.net/pagerank.html>